

Stay in the moment. We handle the after.

• Vultr

• Google · Gemini · ADK

• Speechmatics

A call ends in 30 seconds. The after takes the day.

THE CALL

30s

The booking, in words.
Two humans, done.
Works just fine.

THE AFTER

the rest

Booking to enter.
Confirmation to send.
Allergy to remember.
Briefing for next time.


Every restaurant, dentist, body shop, dog groomer, garage on earth has this gap. Today it's post-its, short-term memory, and a second tab open *just in case*.

**We don't replace
the call.**

**We replace
the after.**

AI receptionists swap out the human. But the human is exactly who customers called for. We keep the human on the phone and rebuild what comes next.

The after is an agentic loop, not a script.



finalize → completed **max_turns → needs_review** **error → failed** no-raise contract — executed actions survive loop errors

The model chooses tools, corrects itself, and exits explicitly.

1 · Self-correction on tool errors

Turn N: `booking.create` → `validation_failed` *party_size required*.

Turn N+1: re-read transcript, re-emit with `party_size=4` → `executed`.

Hard cap: 2 attempts per action.

2 · RAG-on-demand, not RAG-as-prefix

The agent decides whether to spend tokens on the Vector Store — and asks **specific** questions ("allergies on file?"), not catch-all queries. First-time caller: skip entirely.

3 · Explicit loop exit conditions

`finalize` → **completed**. `max_turns` → **needs_review** (first-class terminal, surfaced in the UI). `error` → **failed**. No silent fallback.

4 · Auditable turn-by-turn

Each turn → one audit row. Action results joined on `payload.agent_turn` (integer, not timestamps). The Agent Reasoning Trail is a first-class UI surface, not a buried log.

Real APIs. Billed. Verifiable. One audit trail.

- Vultr

5 services. One audit trail.

`/admin/rag-probe` returns real `input_tokens` — billed on every call.

- **Cloud Compute** · the VM hosting backend + app + demo
- **Coolify on Vultr** · auto-deploy on every push to `main`
- **Managed Postgres 16** · system of record, every audit row
- **Vector Store** · preseeded per-customer memory collection
- **Serverless Inference RAG** · MiniMax-M2.7

verify: `/api/v1/admin/rag-stats`

- Google · Gemini

One model, four agent surfaces.

Every call's Agent Reasoning Trail is one Gemini turn per row — typed tool calls only.

- **gemini-3.1-flash-lite** via Google **ADK 1.18**
- **Call agent** · multi-turn loop, max 12 turns, hard-cap retries
- **Wizard chat** · builds a vertical template in 2–5 messages
- **Briefing regen** · single-shot rewrite of next-call summary
- Structured output: **dynamic Pydantic** ← per-template JSON Schema

verify: any call detail → Agent Reasoning Trail

- Speechmatics

Voice in. Voice out.

Every transcript shows **S1:** / **S2:** — and the agent re-reads spans by speaker.

- **Batch STT** · diarization always on · language=auto
- **Speaker-aware tools** · `search_transcript` + `read_transcript_segment`
- **TTS Preview** · Sonia · Sarah · Megan · Theo voices
- 6 bundled demo MP3s · 2 fresh per wizard-built template
- Bonuses claimed: **multilingual** ✓ · **creative diarization** ✓

verify: any call detail → transcript with S1 / S2

Worldwide problem. Italy is where we measured first.

ITALIAN BASELINE · FLOOR, NOT CEILING

478k

booking-led IT businesses

185k

phone-led subset (SAM)

€110M

initial IT SAM / year
€50 / seat / month

SIZED BY SECTOR First-party Italian sources: FIPE · ISTAT · FNOMCeO · Key-Stone · Confartigianato · AGCOM · UPB.

BEACHHEAD VERTICALS Hair salons · dog groomers · auto repair · beauty centres · dentistry. ~58k restaurants as volume tail.

EXPANSION NEXT France · Germany · UK · Spain · US service economy. Same product, same wizard, single-tenant per customer.

REVENUE MODEL SaaS per seat · SMB owner/operator buyer · expansion by vertical template

Where we differ

vs CallRail · Aircall · Dialpad AI — the AI-receptionist crowd.

	THEM	AFTERGLOW
Where the AI lives	On the call	After the call
Who picks up	AI voice agent	The human
Memory across calls	Transcript log	Briefing + RAG
Agent reasoning	—	First-class UI
Vertical adaptation	Reseller config	2-min wizard
Geo model	Country-specific	Single-tenant · worldwide

Real where it matters. Mocked where a multi-judge demo can't safely write.

- **REAL · billed on every call**

- 1 **Speechmatics**

Batch STT on every uploaded MP3 · diarization on · language auto-detect · paid per minute.

- 2 **Gemini · Google ADK**

Multi-turn agent loop on every analyzed call · paid per token · no offline stub.

- 3 **Vultr RAG**

Non-zero `input_tokens` per query — verifiable live via `/api/v1/admin/rag-probe`.

Also real: Vultr Cloud Compute · Vultr Vector Store (preseeded + write-back in prod) · Vultr Managed Postgres · `customer.update_profile` mutates the row, undoable · Speechmatics TTS Preview generates every bundled and wizard MP3.

- **MOCKED · by design, not by absence**

The outbound write-side integrations no public, multi-judge demo can safely fire. They have a typed contract; we just don't execute them against third parties during the hackathon.

booking.*

whatsapp.*

sms.*

email.*

calendar.*

payment.*

crm.*

review.*

RAG write-back (demo only)

Swap a mock for real = one entry in `action_catalog.py` + an env var. Not an architecture migration. The agent loop, the RAG, and the single-tenant deployment are production-ready today.

Three live URLs. Two-minute click-path to the agent loop.





Live URLs

DEMO (IFRAME)	afterglow.cleversoft.it
OPERATOR APP	app.afterglow.cleversoft.it
API	api.afterglow.cleversoft.it
REPO (MIT)	github.com/Cleversoft-IT/afterglow

Click-path · ~2 minutes

- 1 Open the demo URL. The iframe shows the operator app in a Pixel-style phone frame.
- 2 **Drawer** → **Test simulator** → **Call from existing customer**. Audio plays. Hang up when you've heard enough.
- 3 Back on Home, watch the call: **transcribing** → **analyzing** → **completed**.
- 4 Tap the call. The **Agent Reasoning Trail** on the right shows every turn the agent took.
- 5 Tap "Open contact". The AI-written **briefing** sits at the top of the customer card.
- 6 **Drawer** → **Templates** → **"Set up your own business"**. A 2-minute chat ships a working template — fields, action tools, two fresh demo MP3s.

Multi-visitor sandboxed · per-tab **X-Demo-Session** isolates concurrent judges. **Drawer** → **Settings** → **Reset demo** wipes your sandbox.



afterglow

Stay in the moment.
We handle the after.

`afterglow.cleversoft.it`

`github.com/Cleversoft-IT/afterglow · MIT`

• Vultr

• Google · Gemini · ADK

• Speechmatics

by **ClaudiOpuscoli**