

Stay in the moment.
We handle the after.

• Vultr

• Google · Gemini · ADK

• Speechmatics

A call ends. Something else begins.

THE CALL

Thirty seconds.
Two humans.
Works just fine.

THE AFTER

Booking to enter.
Confirmation to send.
Allergy to remember.
Briefing for next time.


Every restaurant, dentist, body shop, dog groomer, garage on earth has this gap. Today it's post-its, short-term memory, a second tab open *just in case*.

**We don't replace
the call.**

**We replace
the after.**

AI receptionists kill the moat. The reason a customer dials a small business is the human on the other end.

The after is an agentic loop, not a script.



finalize → completed **max_turns** → needs_review **error** → failed no-raise contract — executed actions survive loop errors

The model chooses tools, corrects itself, and exits explicitly.

1 · Self-correction on tool errors

Turn N: `booking.create` → `validation_failed` *party_size required*.

Turn N+1: re-read transcript, re-emit with `party_size=4` → `executed`.

Hard cap: 2 attempts per action.

2 · RAG-on-demand, not RAG-as-prefix

The agent decides whether to spend tokens on the Vector Store — and asks **specific** questions ("allergies on file?"), not catch-all queries. First-time caller: skip entirely.

3 · Explicit loop exit conditions

`finalize` → **completed**. `max_turns` → **needs_review** (first-class terminal, surfaced in the UI). `error` → **failed**. No silent fallback.

4 · Auditable turn-by-turn

Each turn → one audit row. Action results joined on `payload.agent_turn` (integer, not timestamps). The Agent Reasoning Trail is a first-class UI surface, not a buried log.

Real APIs. Billed. Verifiable. One audit trail.

- Vultr

**4 surfaces.
One audit trail.**

`/admin/rag-probe` returns real
`input_tokens` — billed.

Compute + Coolify · Managed Postgres 16 · Vector Store ·
Serverless Inference RAG (MiniMax-M2.7)

- Google · Gemini

**Function-calling agent loop on
ADK.**

Every call's Agent Reasoning Trail is one
Gemini turn per row.

gemini-3.1-flash-lite · ADK 1.18 · typed tools · structured
output: dynamic Pydantic ← JSON Schema

- Speechmatics

**Diarization fed back into the
agent.**

Every transcript shows **S1:** / **S2:** — and the
agent re-reads spans by speaker.

Batch STT · language=auto · TTS Preview for every demo
MP3 · bonuses: multilingual ✓ · creative diarization ✓

Worldwide problem. Italy is where we measured first.

ITALIAN BASELINE · FLOOR, NOT CEILING

478k

booking-led IT businesses

185k

phone-led subset (SAM)

€110M

initial IT SAM / year (€50 / seat)

Sized sector by sector from first-party Italian sources (FIPE, ISTAT, FNOMCeO, Key-Stone, Confartigianato, AGCOM, UPB). Same product, same wizard, expansion next: **France · Germany · UK · Spain · US service economy.**

Beachhead verticals: hair salons · dog groomers · auto repair · beauty · dentistry.

USP vs CallRail · Aircall · Dialpad AI

	THEM	AFTERGLOW
Where the AI lives	On the call	After the call
Who picks up	AI voice agent	The human
Memory across calls	Transcript log	Briefing + RAG
Auditable agent reasoning	—	First-class UI
Vertical adaptation	Reseller config	2-min wizard
Geo model	Country-specific	Single-tenant per customer · worldwide

REVENUE MODEL

SaaS per seat · SMB owner/operator buyer · expansion by vertical template

Real where it matters. Mocked where a multi-judge demo can't safely write.

- **REAL · billed, not stubbed**

- 1 Speechmatics**
Batch STT on every call · diarization on · language=auto.
- 2 Gemini · ADK**
Multi-turn agent loop on every analyzed call. No offline stub.
- 3 Vultr RAG**
Non-zero `input_tokens` per query — verifiable via `/admin/rag-probe` .

Plus Vultr Vector Store · Managed Postgres · `customer.update_profile` row mutation · TTS
Preview for every demo MP3.

- **MOCKED · by design**

The outbound write-side integrations a multi-judge public demo cannot safely fire.

`booking.*` `whatsapp.*` `sms.*` `email.*` `calendar.*` `payment.*`
`crm.*` `review.*` `RAG write-back in demo`

Swap a mock for real = one entry in `action_catalog.py` + an env var. Not an architecture migration.

Three live URLs. Two-minute click-path to the agent loop.

Live URLs

DEMO (IFRAME)	afterglow.cleversoft.it
OPERATOR APP	app.afterglow.cleversoft.it
API	api.afterglow.cleversoft.it
REPO (MIT)	github.com/Cleversoft-IT/afterglow

Click-path · ~2 minutes

- 1 Open the demo URL. Iframe shows the operator app in a phone frame.
- 2 **Drawer** → **Test simulator** → **Call from existing customer**. Audio plays. Hang up.
- 3 Watch the call: **transcribing** → **analyzing** → **completed**.
- 4 Tap the call. Expand the **Agent Reasoning Trail** on the right.
- 5 Tap the customer name. Read the AI-written **briefing** at the top of the card.
- 6 **Drawer** → **Templates** → **Wizard**. Build a dog groomer template in 2 messages.

multi-visitor sandboxed · per-tab **X-Demo-Session** isolates judges

afterglow

Stay in the moment.
We handle the after.

`afterglow.cleversoft.it`

`github.com/Cleversoft-IT/afterglow · MIT`

• Vultr

• Google · Gemini · ADK

• Speechmatics